



Flight Price Prediction

SUBMITTED BY:

SHRAVANI NATAKALA

Introduction

- ▶ Anyone who has booked a flight ticket knows how unexpectedly the prices vary. The cheapest available ticket on a given flight gets more and less expensive over time
- ▶ In first phase we have to collect data of flights ticket from online websites. Here data is collected from “www.yatra.com” website using Selenium technique.
- ▶ Our goal is to build a regression model to predict price of flight ticket.
- ▶ We have also performed the EDA to gain insights of the data.

Data Set Description

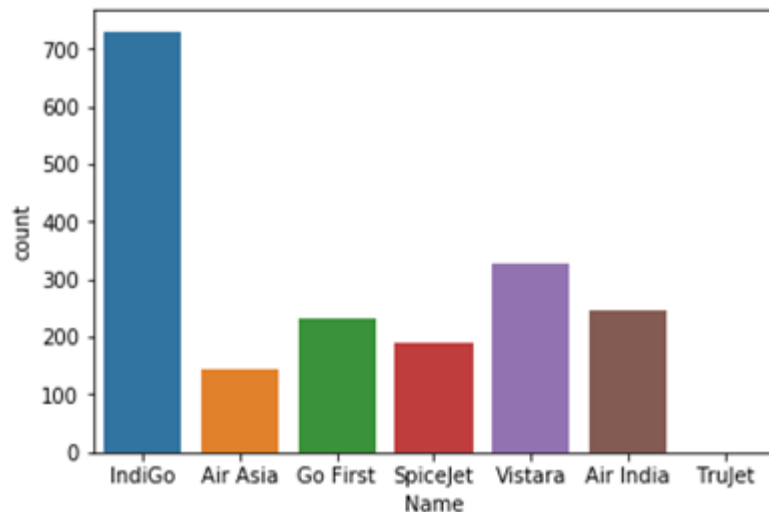
Features:

- ▶ **Name:** name of Airline
- ▶ **Date:** date of journey
- ▶ **Departure:** time of departure
- ▶ **Arrival:** time of arrival
- ▶ **Source:** the source from which service begins
- ▶ **Destination:** the destination where service ends
- ▶ **Stops:** total number of stops between source and destination
- ▶ **Duration:** total duration of flight
- ▶ **Price:** Price of flight ticket

The dataset has no null values.

Data Visualizations

Airline Names




Matrix showing airlines and flights with number of stops


```
#stats of airlines with total stops  
pd.crosstab(df['Name'],df['Stops'])
```

Stops	1 Stop	2 Stop(s)	3 Stop(s)	Non Stop
Name				
Air Asia	86	0	0	58
Air India	147	58	1	42
Go First	129	1	0	101
IndiGo	452	15	0	264
SpiceJet	104	2	0	85
TruJet	0	0	0	1
Vistara	192	32	4	101


Date v/s Price




Name v/s Price



Source v/s Price



Stops v/s Price



Data Pre-Processing

```
# Changing data type of price
p=[]
price=df["Price"]
for i in range(len(price)):
    st=price[i]
    p.append(st.replace(",",""))
df["Price"]=p

df_type_dict={'Price':int}
df=df.astype(df_type_dict)
df.dtypes
```

```
Unnamed: 0      int64
Name           object
Date           object
Departure      object
Arrival        object
Source         object
Destination    object
Stops          object
Duration       object
Price          int32
dtype: object
```

- First we will clean price column by removing ',' and changing it's data type to 'int'

- Next we have removed unnecessary columns and cleaned data in “Arrival”, “Departure”, “Duration” and “Date” and derived new features from each given feature. I have first formed a new data frame and then done all the processing.

```
list=[]
hrs=[]
min=[]
list=time["Departure"]
for i in range(len(list)):
    str=[]
    str=list[i].split(':')
    hrs.append(str[0])    ## Separating hours and minutes
    min.append(str[1])
time["Dep_time_hours"]=hrs
time["Dep_time_min"]=min
```

```
list=[]
hrs=[]
min=[]
list=time["Duration"]
for i in range(len(list)):
    str=[]
    str=list[i].split(' ')
    if(len(str)>1):
        hrs.append(str[0][:-1])    ## Separating hours and minutes
        min.append(str[1][:-1])
    else:
        hrs.append(list[i][:-1])
        min.append('0')
time["Duration_hours"]=hrs
time["Duration_min"]=min
```

```
list=[]
hrs=[]
min=[]
list=time["Arrival"]
for i in range(len(list)):
    str=[]
    str=list[i].split(':')
    hrs.append(str[0])    ## Separating hours and minutes
    min.append(str[1][:3])
time["Arrival_time_hours"]=hrs
time["Arrival_time_min"]=min
```

```
list=[]
d=[]
m=[]
list=time["Date"]
for i in range(len(list)):
    d.append(list[i][:-2])    ## Separating date and month
    m.append(list[i][3:])
time["day"]=d
time["month"]=m
```


- After executing the above lines of code we will get 8 new columns
Dep_time_hours,
Dep_time_min,
Duration_hours,
Duration_min,
Arrival_time_hours,
Arrival_time_min, day and month. Each feature now has integer data type. Since all the usefull information is now extracted we can drop previous columns.

	Departure	Arrival	Duration	Date	Dep_time_hours	Dep_time_min	Duration_hours	Duration_min	Arrival_time_hours	Arrival_time_min	day	month
0	22:40	01:30 + 1 day	2h 50m	24-Jan	22	40	2	50	01	30	24	1
1	3:00	5:55	2h 55m	24-Jan	3	00	2	55	5	55	24	1
2	8:35	11:30	2h 55m	24-Jan	8	35	2	55	11	30	24	1
3	20:15	23:10	2h 55m	24-Jan	20	15	2	55	23	10	24	1
4	9:50	12:50	3h 00m	24-Jan	9	50	3	00	12	50	24	1

- ▶ Next we have introduced two more columns as “Number_of_days” giving the ticket price number of days before the flight service and “Total_duration” giving total time of service.

```
data["month"].value_counts()
```

```
11    744
1     703
12    428
Name: month, dtype: int64
```

```
for i in data["month"]:
    if i==1:
        data["Number_of_days"]=(data["day"]+31+4+30)
    elif i==12:
        data["Number_of_days"]=(data["day"]+30+4)
    elif i==11:
        data["Number_of_days"]=(data["day"]+4)
    else:
        print("Add another condition")
```

```
#introducing a new column total duration
data["Total_duration"]=data["Duration_hours"]*60+data["Duration_min"]
```

- ▶ Encoding variables with object data type: We have encoded “Stops” manually and used LabelEncoder for other variables.
- ▶ We also observed outliers and skewness in data for which we used z-score method and log transformation to deal with it. In this process we faces a data loss of 2.5%.

```
# Encoding Total_Stops Column
data["Stops"] = data["Stops"].replace({'Non Stop':0, '1 Stop':1, '2 Stop(s)':2, '3 Stop(s)':3, '4 Stop(s)':4})
data.head()
```

```
lab_enc=LabelEncoder()
cols=["Name", "Source", "Destination"]
for i in cols:
    df1= lab_enc.fit_transform(data[i])
    data[i]=df1
data.head()
```

PREPARING DATA FOR MODEL

- ▶ Making our Data ready for model Building phase we will first separate target variable from other features. Then use `StandardScaler` to scale data and use `train_test_split` to split data into train and test to make it ready for model.
- ▶ For `train_test_split` we found the best random state by running a loop on linear regression and checking for best accuracy.

MODEL BUILDING AND EVALUATION

Algorithms used are:

- ▶ Linear Regression
- ▶ Decision Tree Regressor
- ▶ KNN Regressor
- ▶ Random Forest Regressor
- ▶ Gradient Boosting Regressor

Linear Regression

```
**** LinearRegression ****
```

```
accuracy_score: 0.5852425650963049
```

```
cross_val_score: 0.5727340366915723
```

```
mean_squared_error 2423438.683034735
```

KNeighboursRegressor

```
**** KNeighborsRegressor ****
```

```
accuracy_score: 0.7460475895421008
```

```
cross_val_score: 0.7698852848280546
```

```
mean_squared_error 1483850.6639344261
```

DecisionTreeRegressor

```
**** DecisionTreeRegressor ****
```

```
accuracy_score: 0.7599649123898086
```

```
cross_val_score: 0.8201500680455596
```

```
mean_squared_error 1402531.3777322404
```

RandomForestRegressor

```
**** RandomForestRegressor ****
```

```
accuracy_score: 0.8606367626369499
```

```
cross_val_score: 0.8886455523278205
```

```
mean_squared_error 814303.0889777505
```



GradientBoostingRegressor

```
**** GradientBoostingRegressor ****
```

```
accuracy_score: 0.8261236774436168
```

```
cross_val_score: 0.8435510487665692
```

```
mean_squared_error 1015963.9603442102
```


Choosing Best Model

After running the loop we get a dataframe showing each model and scores obtained from it.

Looking the various metrics we conclude “**Random Forest Model**” as our best model and hence we will now tune our model.

	Model	Accuracy_score	Cross_val_score	Mean_Squared_Error
0	LinearRegression	58.524257	57.273404	2.423439e+06
1	KNeighborsRegressor	74.604759	76.988528	1.483851e+06
2	DecisionTreeRegressor	75.996491	82.015007	1.402531e+06
3	RandomForestRegressor	86.063676	88.864555	8.143031e+05
4	GradientBoostingRegressor	82.612368	84.355105	1.015964e+06

Hyper-Parametric Tuning

```
In [63]: rmf= RandomForestRegressor()
params={'max_features':['auto','sqrt'],'n_estimators':[50,80], 'criterion':['mse','mae'],
        'max_depth':[5,10], 'min_samples_split':[4,6],
        'min_samples_leaf':[2,3]}
grd=GridSearchCV(rmf,param_grid=params)
grd.fit(x_train,y_train)
print('best params=>',grd.best_params_)

best params=> {'criterion': 'mse', 'max_depth': 10, 'max_features': 'auto', 'min_samples_leaf': 2, 'min_samples_split': 4, 'n_estimators': 50}
```

```
In [64]: rmf= RandomForestRegressor(criterion= 'mse', max_depth= 10, max_features= 'auto', min_samples_leaf= 2,
min_samples_split= 4, n_estimators= 50)
rmf.fit(x_train,y_train)
y_pred=rmf.predict(x_test)
print("Random Forest Regression: Accuracy = ",rmf.score(x_test,y_test))
print("\n Mean Squared Error= ",mean_squared_error(y_test,y_pred))
print("\n Root Mean Squared Error= ",np.sqrt(mean_squared_error(y_test,y_pred)))
print("\n Mean Absolute Error= ",mean_absolute_error(y_test,y_pred))
```

Random Forest Regression: Accuracy = 0.8453601476367489

Mean Squared Error= 903564.7552476081


Root Mean Squared Error= 950.5602323091409

Mean Absolute Error= 547.8271263504754

Final prediction gives 84.53% accuracy

Conclusion

- ▶ First, we collected data on flight ticket prices from “yatra.com”, it was done by using Web scraping. The framework used for web scraping was Selenium.
- ▶ Then the scrapped data was saved in a csv file to use it for modeling purpose.
- ▶ From the extensive EDA performed in this project we observed:
 - a) Flights from Bangalore and Jammu have higher prices.
 - b) Flights with longer route i.e. high number of stops have high prices.
 - c) Also, prices of flight in next month are high as compared to those in coming months.
 - d) From the given data we can also conclude that AirIndia and vistara flights are expensive as compared to other flights.
- ▶ The model build after hyper-parametric tuning gives an accuracy for 84.53%

- 
- ▶ After the completion of this project, we got an insight of how to collect data, pre-processing the data, analyzing the data and building a model. It helped me to gain conclusions from graphs. Also it helped me in exploring multiple algorithms and metrics to get the best output.
 - ▶ Since the data keeps changing we cannot fully rely on this project in the distant future we need to update it with updation in data
 - ▶ Also the scrapping of data took a lot of time as there was no such detail mentioned on fetching data. Random sources and destinations are used to pick up data.
 - ▶ This project is done with limited resources and can be made more efficient in future..